

# Detecting Vulnerabilities in Virtual Devices with Conformance Testing

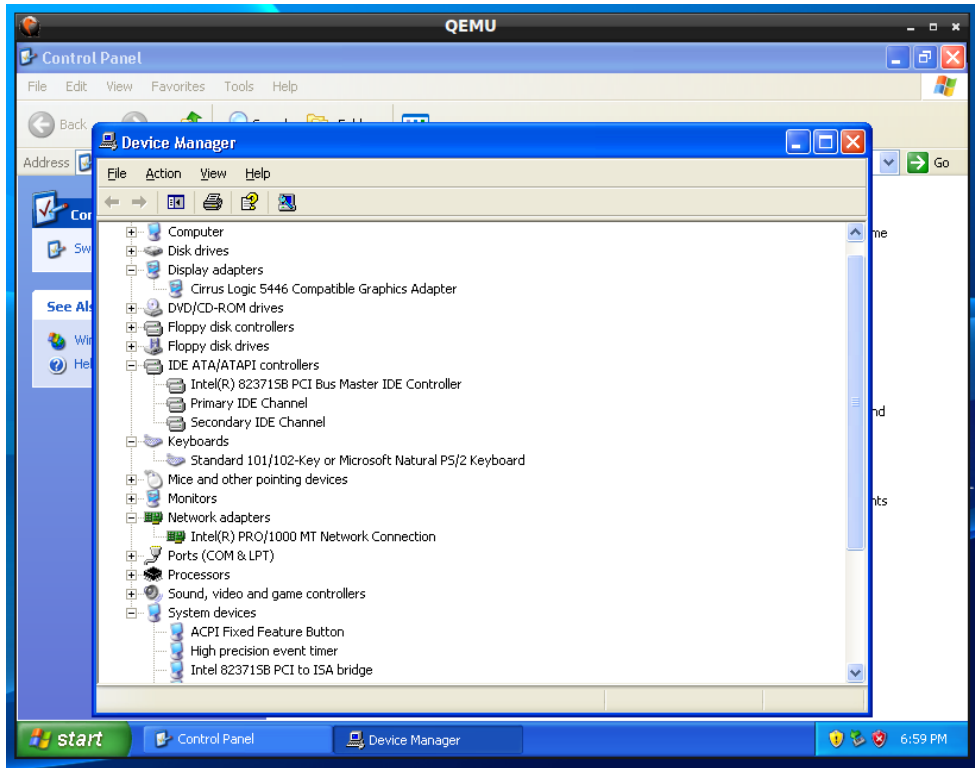
-- Turning old hardware Into gold



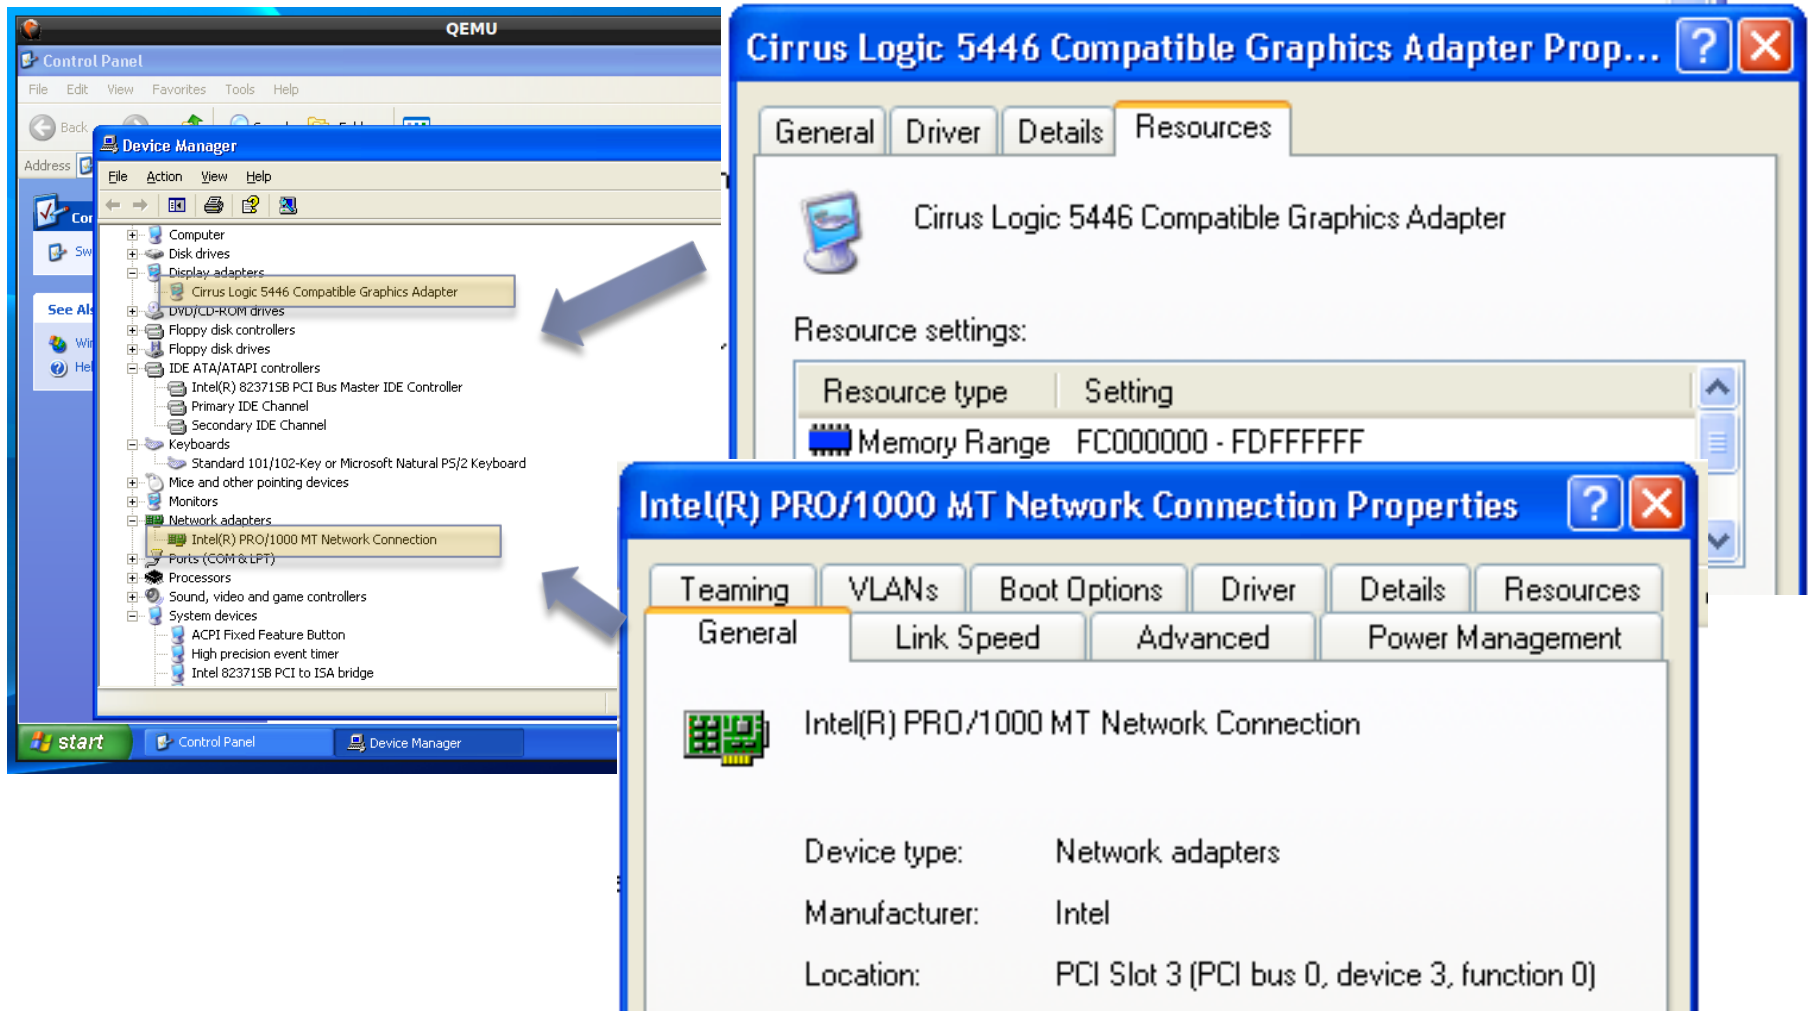
Kang Li  
Michael Contreras

**kodos @disekt.org**  
**maikol@disekt.org**

# VM and Virtual Devices



# QEMU and its Virtual Devices



# QEMU and its Virtual Devices

---

## ▶ QEMU

- ▶ Open-source hypervisor with hardware virtualization
- ▶ Share common components with KVM and VirtualBox

## ▶ Many types of devices

- ▶ Graphics
- ▶ Network
- ▶ Disk Controller
- ▶ USB
- ▶ ...



# QEMU and its Virtual Devices

---

- ▶ Multiple devices in the same category
  - ▶ Network interfaces



Intel  
E1000




Intel  
eepro100



Novell  
NE2000



AMD  
PC-NET II



Realtek  
RTL8139



SMCS  
LAN9118

- ▶ What's common about these devices?
    - ▶ All pictures copied from ebay, all below \$10.
    - ▶ Old devices that you have probably thrown away!
    - ▶ VMs (and Cloud) are “faking” these old devices.
-

# VM and Cloud provide an illusion

---





# When virtual devices are done right

---



# What if the virtual device goes wrong?

---



# Ensuring virtual devices are done right

---

Conformance Testing --

compare a virtual device with its physical counterpart

---

# Conformance Testing is like Comparing ...

---



Original (**OLD**)



Remake (**NEW**)  
/ Mimic



# Example: Mishandled Jumbo Frames

- ▶ We discovered memory corruption bug
  - ▶ Reported on 12/02/2012
  - ▶ Confirmed on 12/06/2012 (CVE 2012-6075)
- ▶ Assumptions about the hardware
  - ▶ Large packet should be dropped when LPE is disabled.
  - ▶ A misunderstanding between driver and device
- ▶ Driver optimizes allocations based on this assumption
  - ▶ Driver allocates 1522 bytes for buffers, expecting anything larger will be dropped by the NIC
  - ▶ Physical hardware respects this limit but virtual device does not
- ▶ Mismatching behavior causes overwrite in Guest OS memory



# Example: Mishandled Jumbo Frames


- ▶ Fuzzed a physical Intel E1000 device
- ▶ Captured some of the device MMIO registers during the fuzz
  - ▶ E.g. GPRC, RLEC
- ▶ Replayed the I/O events (traces) on the QEMU virtual device
- ▶ No immediate crash
- ▶ However, detected the following inconsistencies
  - ▶ **RLEC** @ 0x04040 – Receive Length Error Count
  - ▶ **PRC** @ 0x0405C – Packets Received ([64-1522] Bytes) Count
  - ▶ **BPRC** @ 0x04078 – Broadcast Packets Received Count
  - ▶ **MPRC** @ 0x0407C – Multicast Packets Received Count
  - ▶ **GPRC** @ 0x04074 – Good Packet Received Count



# What can be observed/compared

---

- ▶ Full visibility of virtual devices
- ▶ But limited visibility of physical devices
- ▶ Observed by Capturing Traces



Trace:

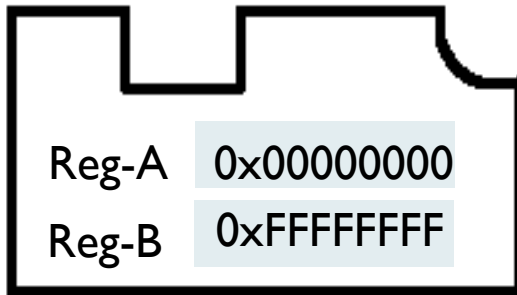
- Event:** mmio\_write(reg, value)
- Device State:**  $[R_1, R_2, \dots, R_N]$
- Event:** mmio\_read(reg)
- Device State:**  $[R_1, R_2, \dots, R_N]$
- ... ..
- Event:** mmio\_write(reg, value)
- Device States:**  $[R_1, R_2, \dots, R_N]$
- ... ..

# Conformance Testing (Ideally)

---

**Spec:** Reg-A is a mask register for Reg-B.

An update to A causes B to change to  $V_B \& \sim V_A$



HW before I/O event

---

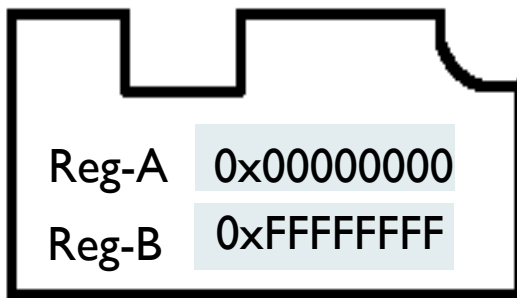


# Conformance Testing (Ideally)



---

**Spec:** Reg-A is a mask register for Reg-B.

An update to A causes B to change to  $V_B \& \sim V_A$



HW before I/O event



HW after I/O event


---

# Conformance Testing (Ideally)



---

**Spec:** Reg-A is a mask register for Reg-B.

An update to A causes B to change to  $V_B \& \sim V_A$



HW before I/O event




HW after I/O event

**Reproduce the above  
operation using the  
virtual device**


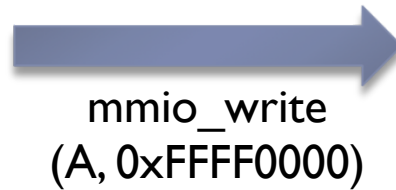
# Conformance Testing (Ideally)

**Spec:** Reg-A is a mask register for Reg-B.

An update to A causes B to change to  $V_B \& \sim V_A$




HW before I/O event





HW after I/O event

**Reproduce the above  
operation using the  
virtual device**



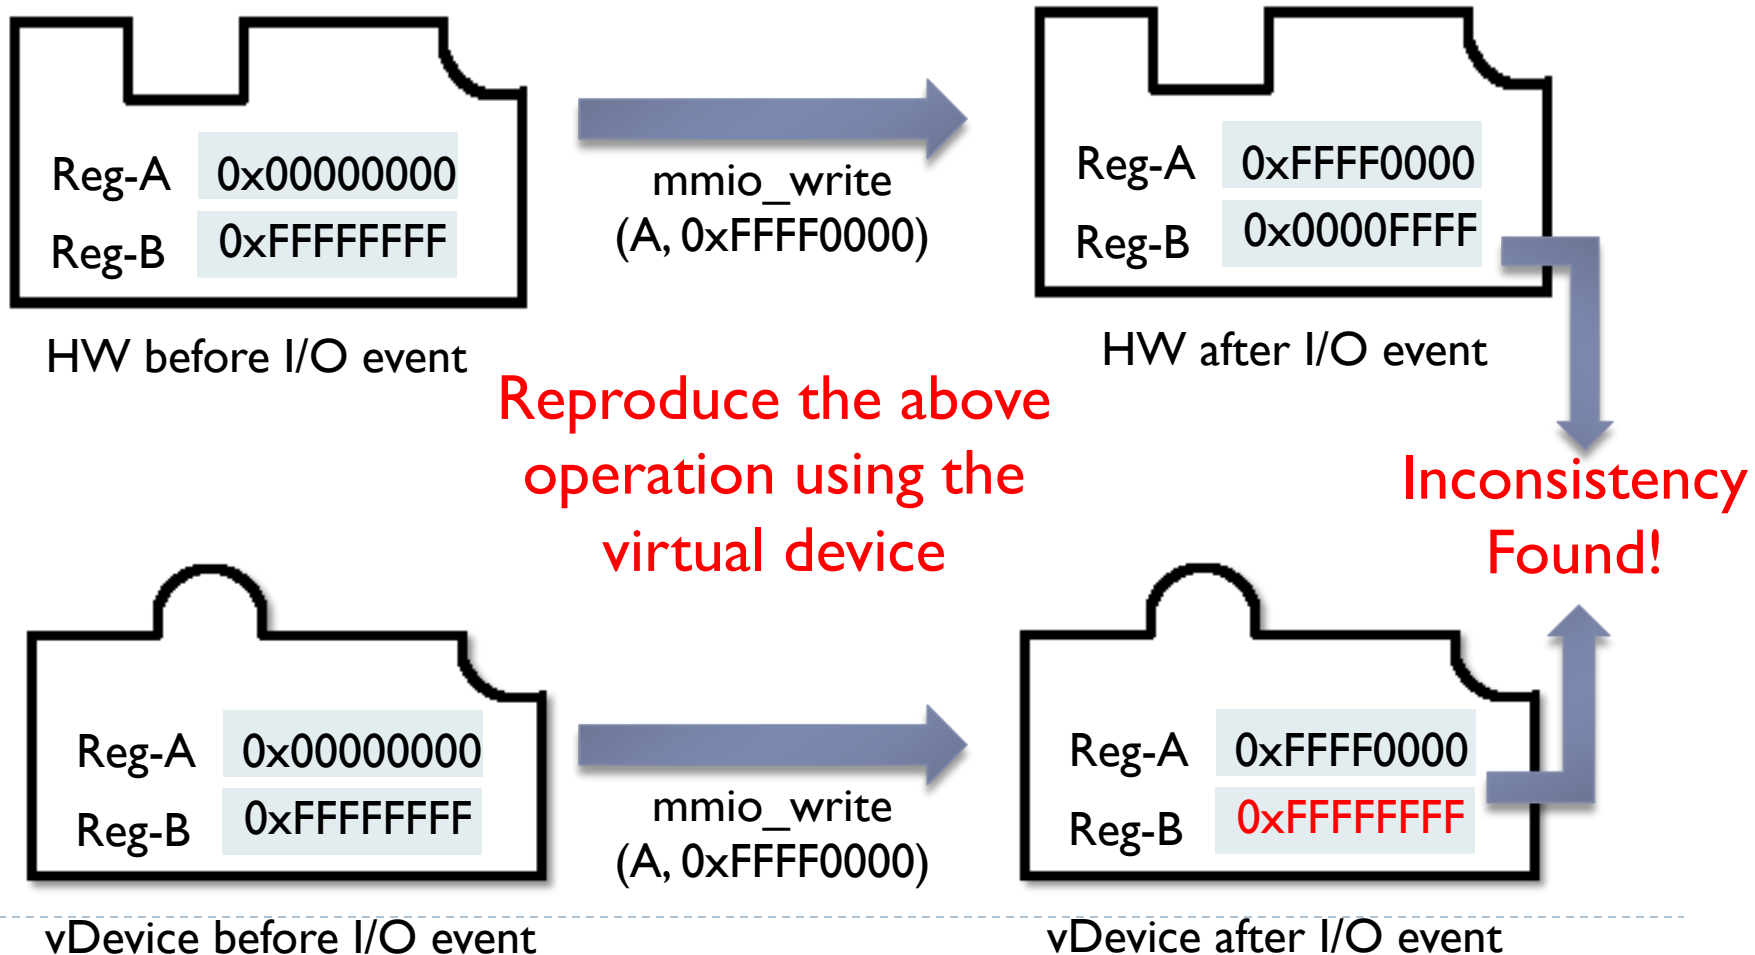
vDevice before I/O event



vDevice after I/O event

# Conformance Testing (Ideally)

**Spec:** Reg-A is a mask register for Reg-B.  
An update to A causes B to change to  $V_B \& \sim V_A$



# Conformance Testing (In Reality)

---

- ▶ Dump and replay only works in simple cases
    - ▶ Not all physical registers are accessible (readable)
    - ▶ Some events are difficult to observe or expensive to capture
      - ▶ E.g. DMA
    - ▶ Some registers are accessible, but have side effects
      - ▶ E.g., clear-on-read
-

# Symbolic Conformance Testing

---



- ▶ Symbolic execution is getting popular for bug finding
    - ▶ KLEE symbolic execution engine
      - ▶ <http://klee.lvm.org/>
      - ▶ OSDI Paper from Cristian Cadar, Daniel Dunbar, and Dawson Engler
    - ▶ Users of symbolic execution
      - ▶ Coverity, ForAllSecure, and bugchecker.net ...
  
  - ▶ Our approach to deal with unobservable device state
    - ▶ Construct the virtual device state by setting
      - ▶ observable register values based on the trace
      - ▶ missing registers with symbolic values
-

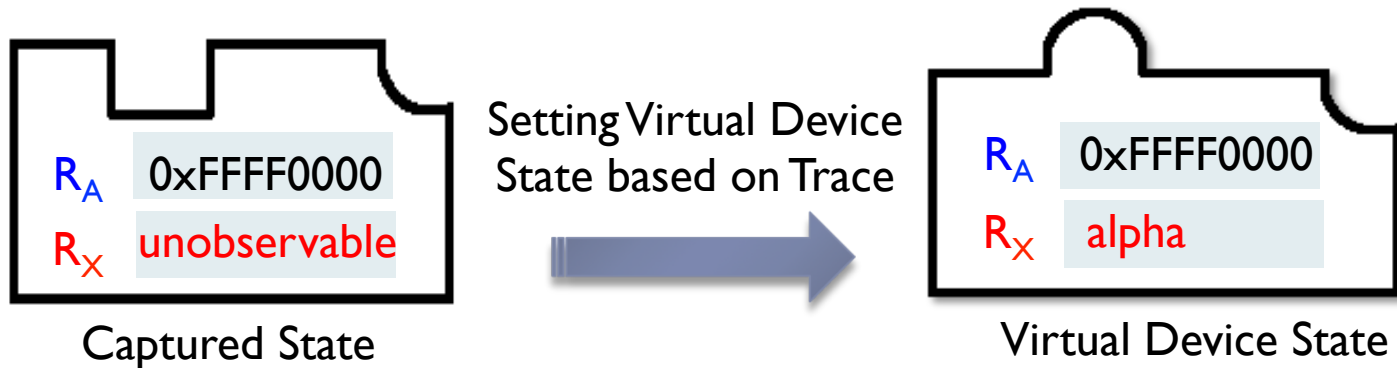


# Symbolic Register Values

---

## ► Example:

- For a dumb device with only 2 registers:
  - $R_A$  (observable) and  $R_X$  (unobservable)
- The device state in a trace looks like this: [ $R_A == 0xFFFF0000$ ]



# How to Run with Symbolic Values?

---

- ▶ Consider the following virtual device program:

## Virtual Device Code Snippet


```
...  
mmio_write_update_RA (value)  
{  
    if ( $R_X == 0$ )  
         $R_A = \text{value};$   
    else  
         $R_A = \text{value} \& 0x0000FFFF;$   
}
```

# How to Run with Symbolic Values?

- ▶ Consider the following virtual device program:

## Virtual Device Code Snippet

```
...  
mmio_write_update_RA (value)  
{  
    if ( $R_X == 0$ )  
         $R_A = \text{value}$ ;  
    else  
         $R_A = \text{value} \& 0x0000FFFF$ ;  
}
```



Virtual Device State

+ **Event I:**  
write ( $R_A$ ,  $0xC0FFEE$ )


**Suppose we have the  
above initial state  
and a given event ...**

# How to Run with Symbolic Values?

- ▶ Consider the following virtual device program:


## Virtual Device Code Snippet

```
...  
mmio_write_update_RA (value)  
{  
    if ( $R_X == 0$ )  
         $R_A = \text{value}$ ;  
    else  
         $R_A = \text{value} \& 0x0000FFFF$ ;  
}
```



Virtual Device State

+ Event I:  
write ( $R_A$ ,  $0xC0FFEE$ )




**What will the virtual device state be after Event I?**

# Symbolic Execution

- ▶ Consider the following virtual device program:

## Virtual Device Code Snippet


```
...  
mmio_write_update_RA (value)  
{  
    if ( $R_X == 0$ )  
         $R_A = \text{value}$ ;  
    else  
         $R_A = \text{value} \& 0x0000FFFF$ ;  
}
```



Virtual Device State

+ **Event I:**  
write ( $R_A$ ,  $0xC0FFEE$ )

If ( $\text{alpha} == 0$ )  
Transaction #1




Virtual Device State

# Symbolic Execution

- Consider the following virtual device program:

## Virtual Device Code Snippet


```
...  
mmio_write_update_RA (value)  
{  
    if ( $R_X == 0$ )  
         $R_A = \text{value}$ ;  
    else  
         $R_A = \text{value} \& 0x0000FFFF$ ;  
}
```



Virtual Device State


+ **Event I:**  
write ( $R_A$ ,  $0xC0FFEE$ )

**If ( $\text{alpha} == 0$ )**  
Transaction #1



Virtual Device State

**If ( $\text{alpha} != 0$ )**  
Transaction #2



Virtual Device State



# Searching for Inconsistencies

**Given this Captured Trace:**

...


Device State: [ $R_A == 0xFFFF0000$ ]

Event I: mmio\_write ( $R_A$ , 0xC0FFEE)

**Device State: [ $R_A == 0xFFEE$ ]**

...


- Does one of the output virtual device states match the captured device state?



Virtual Device State


+ **Event I:**  
write ( $R_A$ , 0xC0FFEE)

**If ( $\alpha == 0$ )**  
Transaction #1



Virtual Device State

**If ( $\alpha != 0$ )**  
Transaction #2



Virtual Device State

# Searching for Inconsistencies

Given this Captured Trace:


...

Device State: [ $R_A == 0xFFFF0000$ ]

Event I: mmio\_write ( $R_A$ , 0xC0FFEE)

Device State: [ $R_A == 0xFFEE$ ]


...



Virtual Device State


+ Event I:  
write ( $R_A$ , 0xC0FFEE)

If ( $\alpha == 0$ )  
Transaction #1



Virtual Device State

If ( $\alpha != 0$ )  
Transaction #2



Virtual Device State

- Found a match, continue with the Transaction.
- If multiple matches found, follow each one.

# Searching for Inconsistencies (cont.)

## Given this Captured Trace:

...

Device State: [ $R_A == 0xFFFF0000$ ]

Event I: mmio\_write ( $R_A$ , 0xC0FFEE)


Device State: [ $R_A == 0xFFEE$ ]

**Event II:** mmio\_write ( $R_A$ , 0x00BEFF00)

**Device State:** [ $R_A == 0xBEFF00$ ]

...

## Follow from previous transaction



Virtual Device State

**Event II:**  
+ write ( $R_A$ , 0x00BEFF00)

- ▶ Checking a trace with consecutive events

# Searching for Inconsistencies (cont.)

## Given this Captured Trace:

...

Device State: [ $R_A == 0xFFFF0000$ ]

Event I: mmio\_write ( $R_A$ , 0xC0FFEE)


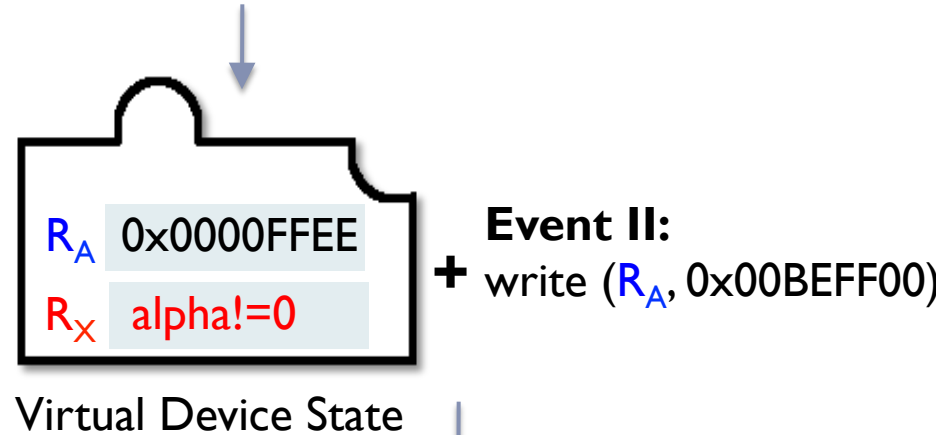
Device State: [ $R_A == 0xFFEE$ ]

**Event II:** mmio\_write ( $R_A$ , 0x00BEFF00)

**Device State:** [ $R_A == 0xBEFF00$ ]

...


## Follow from previous transaction



- ▶ Checking a trace with consecutive events
- ▶ No candidate match → Inconsistency Found!

# Summary

---

- ▶ Conformance Testing 
    - ▶ A unique way to test virtual device implementations.
  - ▶ Symbolic Execution helps to deal with partial states.
  - ▶ QEMU bugs can be detected in this way.
    - ▶ Many inconsistencies found in EEPROM, E1000, and other virtual devices.
-